

Decision Tree and Naive Bayes(1)

November 9, 2021

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
from collections import Counter
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
```

```
[3]: Prudential_train = pd.read_csv("https://raw.githubusercontent.com/crisajose/
↳CIND-820-Big-Data-Analytics-Project/main/train.csv")
```

```
[4]: Prudential_train.head()
```

```
[4]:
```

	Id	Product_Info_1	Product_Info_2	Product_Info_3	Product_Info_4	\
0	2	1	D3	10	0.076923	
1	5	1	A1	26	0.076923	
2	6	1	E1	26	0.076923	
3	7	1	D4	10	0.487179	
4	8	1	D2	26	0.230769	

	Product_Info_5	Product_Info_6	Product_Info_7	Ins_Age	Ht	...	\
0	2	1	1	0.641791	0.581818	...	
1	2	3	1	0.059701	0.600000	...	
2	2	3	1	0.029851	0.745455	...	
3	2	3	1	0.164179	0.672727	...	
4	2	3	1	0.417910	0.654545	...	

	Medical_Keyword_40	Medical_Keyword_41	Medical_Keyword_42	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	Medical_Keyword_43	Medical_Keyword_44	Medical_Keyword_45	\
0	0	0	0	

1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

	Medical_Keyword_46	Medical_Keyword_47	Medical_Keyword_48	Response
0	0	0	0	8
1	0	0	0	4
2	0	0	0	8
3	0	0	0	8
4	0	0	0	8

[5 rows x 128 columns]

```
[5]: CATEGORICAL_COLUMNS = ["Product_Info_1", "Product_Info_2", "Product_Info_3",\
    ↪ "Product_Info_5", "Product_Info_6",\
    ↪ "Product_Info_7", "Employment_Info_2",\
    ↪ "Employment_Info_3", "Employment_Info_5", "InsuredInfo_1",\
    ↪ "InsuredInfo_2", "InsuredInfo_3", "InsuredInfo_4",\
    ↪ "InsuredInfo_5", "InsuredInfo_6", "InsuredInfo_7",\
    ↪ "Insurance_History_1", "Insurance_History_2",\
    ↪ "Insurance_History_3", "Insurance_History_4", "Insurance_History_7",\
    ↪ "Insurance_History_8", "Insurance_History_9",\
    ↪ "Family_Hist_1", "Medical_History_2", "Medical_History_3",\
    ↪ "Medical_History_4", "Medical_History_5",\
    ↪ "Medical_History_6", "Medical_History_7", "Medical_History_8",\
    ↪ "Medical_History_9", "Medical_History_11",\
    ↪ "Medical_History_12", "Medical_History_13", "Medical_History_14",\
    ↪ "Medical_History_16", "Medical_History_17",\
    ↪ "Medical_History_18", "Medical_History_19", "Medical_History_20",\
    ↪ "Medical_History_21", "Medical_History_22",\
    ↪ "Medical_History_23", "Medical_History_25", "Medical_History_26",\
    ↪ "Medical_History_27", "Medical_History_28",\
    ↪ "Medical_History_29", "Medical_History_30", "Medical_History_31",\
    ↪ "Medical_History_33", "Medical_History_34",\
    ↪ "Medical_History_35", "Medical_History_36", "Medical_History_37",\
    ↪ "Medical_History_38", "Medical_History_39",\
    ↪ "Medical_History_40", "Medical_History_41"]
CONTINUOUS_COLUMNS = ["Product_Info_4", "Ins_Age", "Ht", "Wt", "BMI",
    ↪ "Employment_Info_1", "Employment_Info_4",\
    ↪ "Employment_Info_6",
    ↪ "Insurance_History_5", "Family_Hist_2", "Family_Hist_3",\
    ↪ "Family_Hist_4", "Family_Hist_5"]
DISCRETE_COLUMNS = ["Medical_History_1", "Medical_History_10",\
    ↪ "Medical_History_15", "Medical_History_24", "Medical_History_32"]
DUMMY_COLUMNS = ["Medical_Keyword_{i}".format(i) for i in range(1, 48)]
```

```
[6]: categorical_data = Prudential_train[CATEGORICAL_COLUMNS]
```

```
[7]: continuous_data = Prudential_train[CONTINUOUS_COLUMNS]
```

```
[8]: discrete_data = Prudential_train[DISCRETE_COLUMNS]
```

```
[9]: dummy_data = Prudential_train[DUMMY_COLUMNS]
```

1 Variable Types

```
[10]: Prudential_train.dtypes
```

```
[10]: Id                int64
      Product_Info_1    int64
      Product_Info_2    object
      Product_Info_3    int64
      Product_Info_4    float64
      ...
      Medical_Keyword_45 int64
      Medical_Keyword_46 int64
      Medical_Keyword_47 int64
      Medical_Keyword_48 int64
      Response          int64
      Length: 128, dtype: object
```

2 NULL values

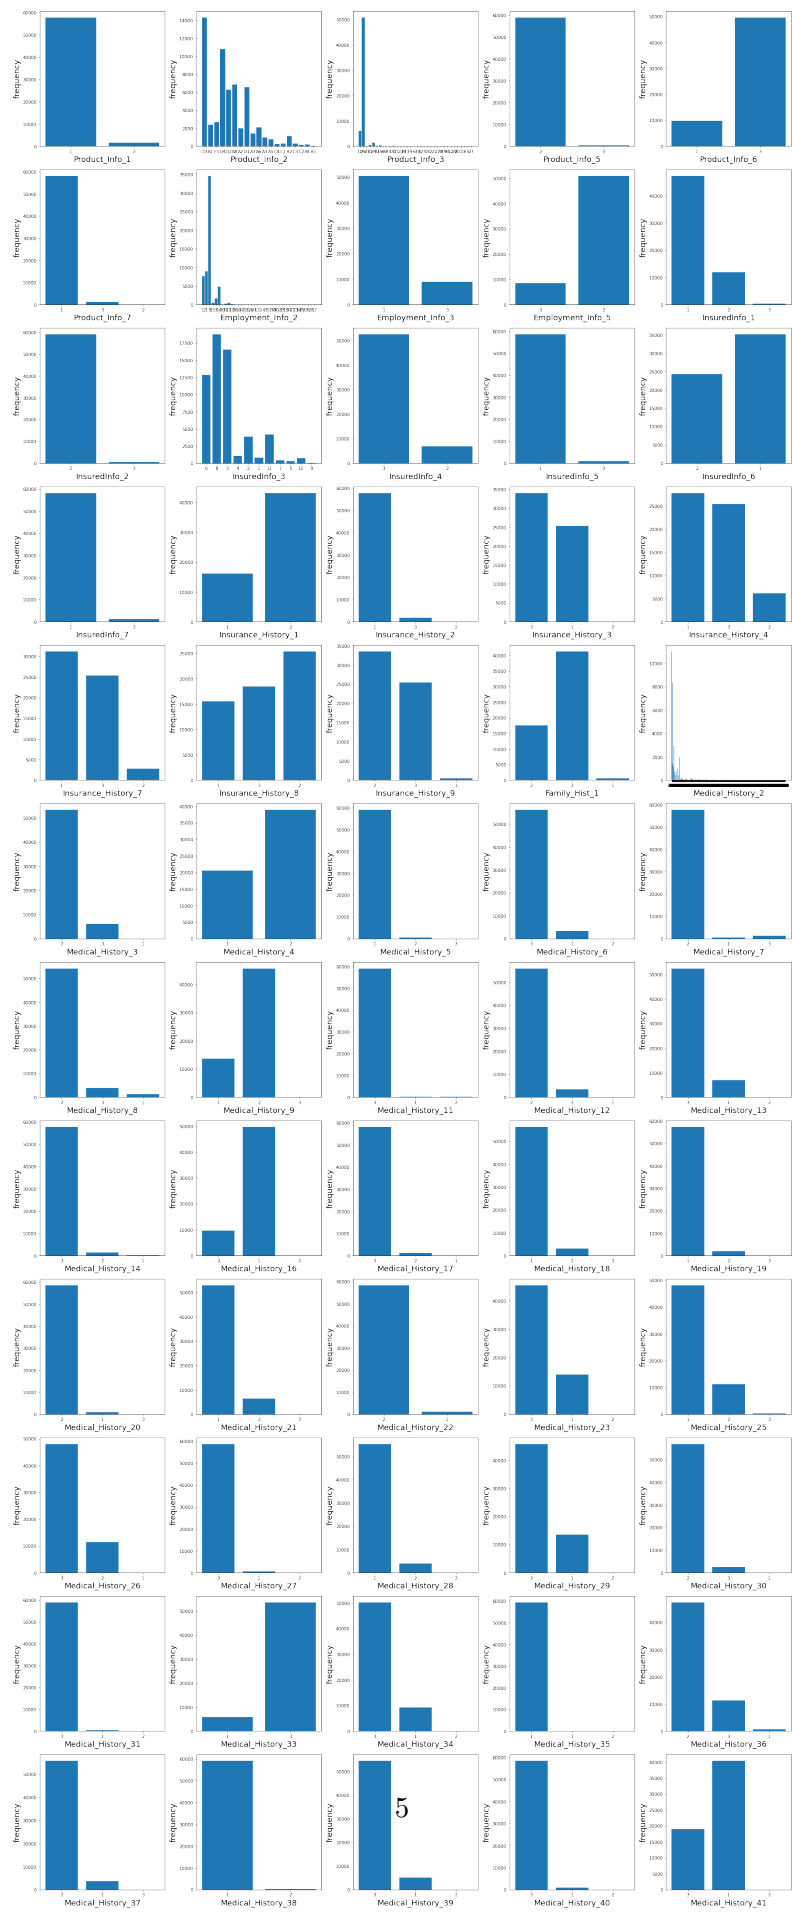
```
[12]: Prudential_NULL = Prudential_train.isnull().any()
      Prudential_NULL = [col for col in Prudential_train.columns if
      ↪ Prudential_train[col].isnull().any()]
      Prudential_NULL
```

```
[12]: ['Employment_Info_1',
      'Employment_Info_4',
      'Employment_Info_6',
      'Insurance_History_5',
      'Family_Hist_2',
      'Family_Hist_3',
      'Family_Hist_4',
      'Family_Hist_5',
      'Medical_History_1',
      'Medical_History_10',
      'Medical_History_15',
```

```
'Medical_History_24',  
'Medical_History_32']
```

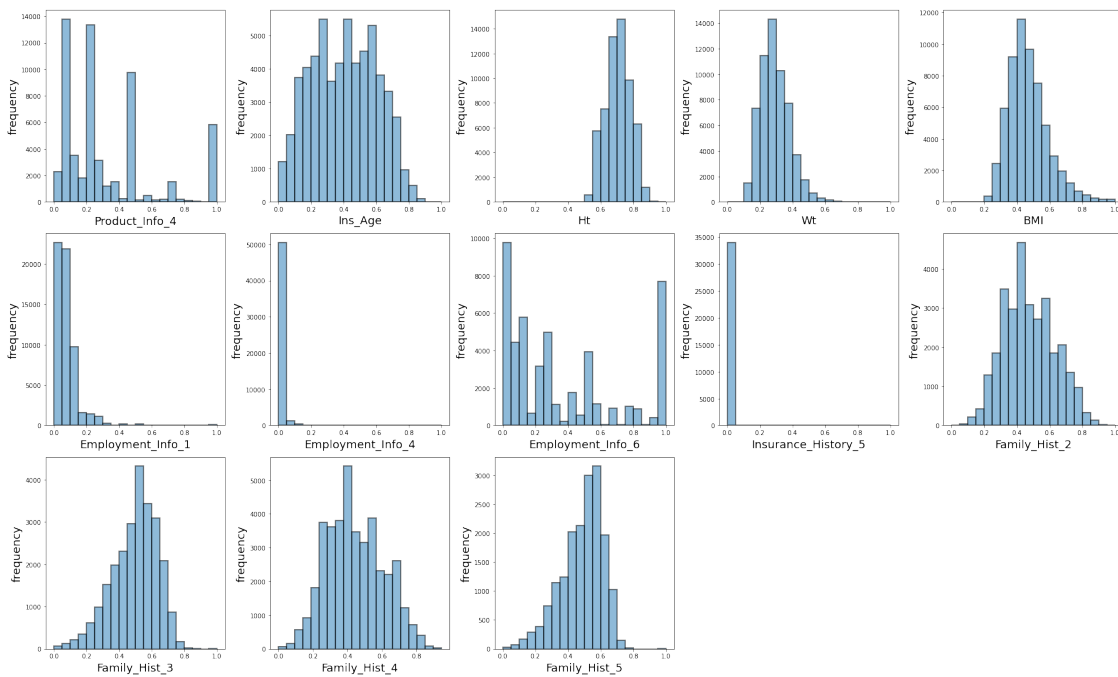
3 Categorical Variable - Plot

```
[13]: def plot_categoricals(data):  
    ncols = len(data.columns)  
    fig = plt.figure(figsize=(5 * 5, 5 * (ncols // 5 + 1)))  
    for i, col in enumerate(data.columns):  
        cnt = Counter(data[col])  
        keys = list(cnt.keys())  
        vals = list(cnt.values())  
        plt.subplot(ncols // 5 + 1, 5, i + 1)  
        plt.bar(range(len(keys)), vals, align="center")  
        plt.xticks(range(len(keys)), keys)  
        plt.xlabel(col, fontsize=18)  
        plt.ylabel("frequency", fontsize=18)  
    fig.tight_layout()  
    plt.show()  
  
plot_categoricals(categorical_data)
```




4 Continuous Variable - Plot

```
[14]: def plot_histgrams(data):
    ncols = len(data.columns)
    fig = plt.figure(figsize=(5 * 5, 5 * (ncols // 5 + 1)))
    for i, col in enumerate(data.columns):
        X = data[col].dropna()
        plt.subplot(ncols // 5 + 1, 5, i + 1)
        plt.hist(X, bins=20, alpha=0.5, \
                 edgecolor="black", linewidth=2.0)
        plt.xlabel(col, fontsize=18)
        plt.ylabel("frequency", fontsize=18)
    fig.tight_layout()
    plt.show()
plot_histgrams(continuous_data)
```



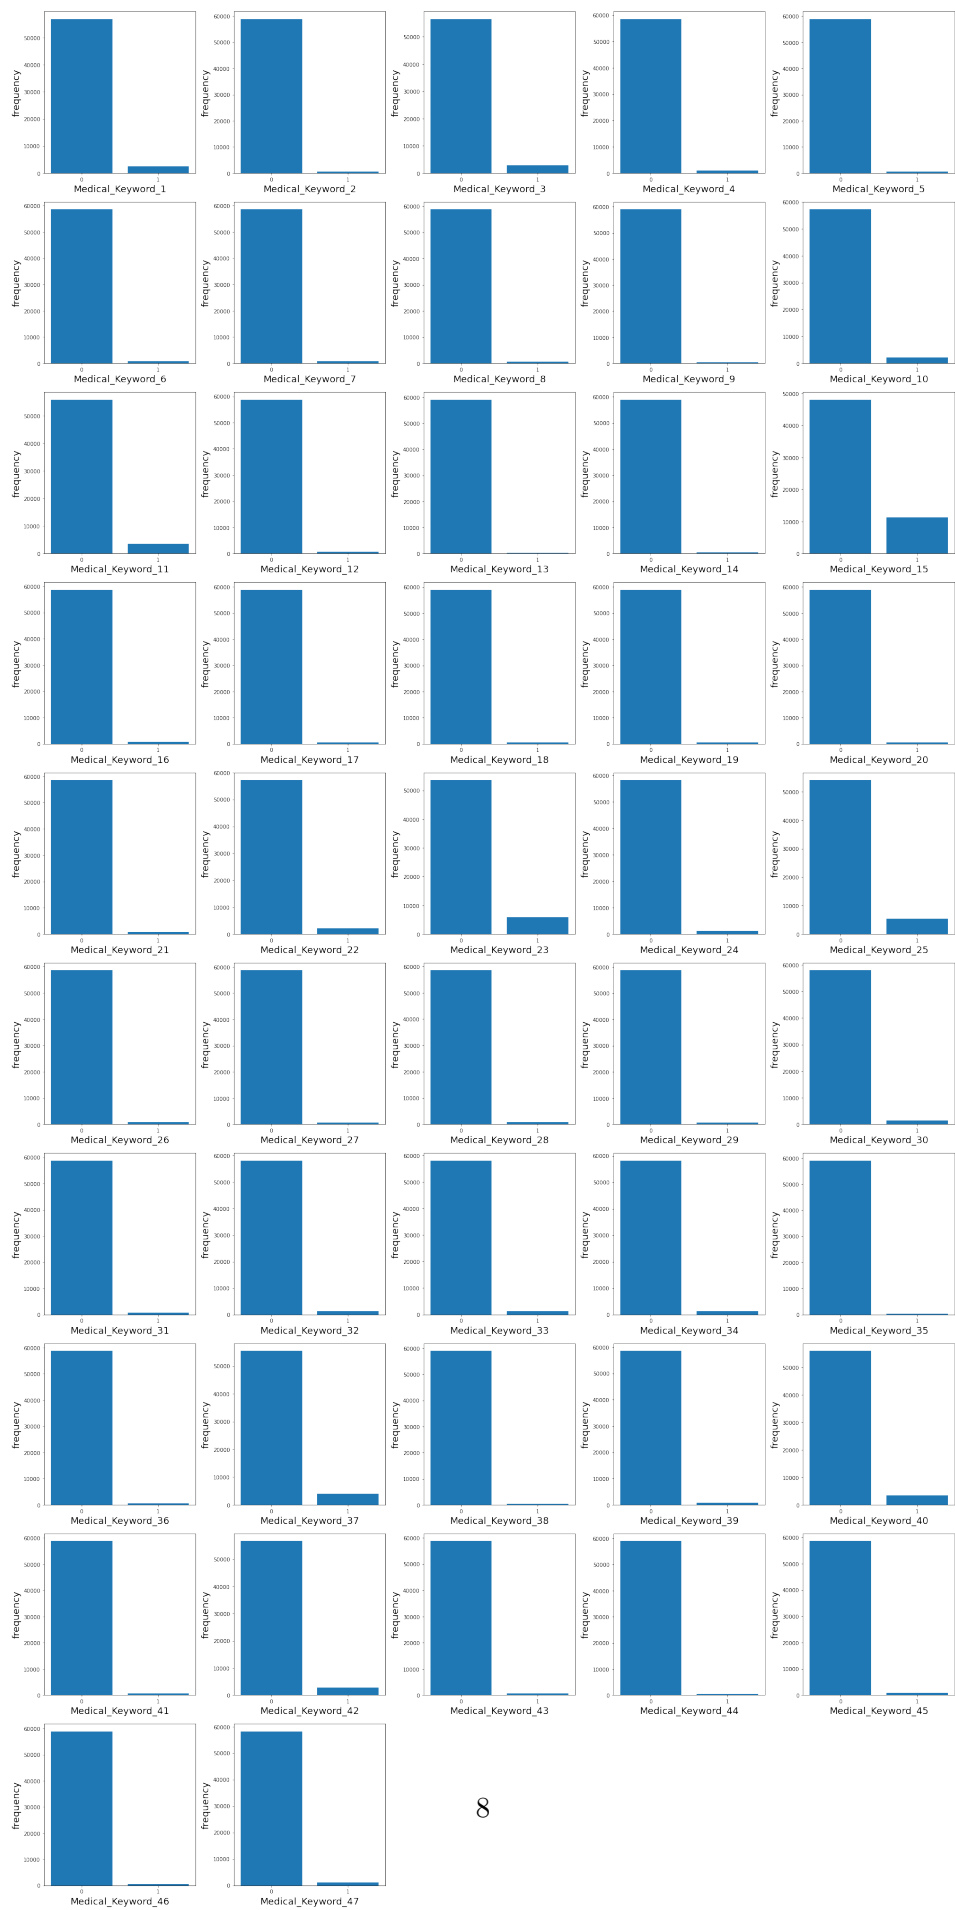
5 Discrete Variable - Plot

```
[16]: plot_histograms(discrete_data)
```



6 Dummy Variable - Plot


```
[17]: plot_categoricals(dummy_data)
```



7 Response Data Distribution

```
[18]: sns.countplot(x=Prudential_train.Response).set_title('Distribution of rows by_  
→response categories')
```


```
[18]: Text(0.5, 1.0, 'Distribution of rows by response categories')
```



8 Outliers Plot

```
[19]: misc_cols=["Ins_Age", "Ht", "Wt", "BMI"]  
sns.boxplot(data=Prudential_train[misc_cols])
```

```
[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa0414a8ed0>
```



9 Reassign Risk Class

```
[20]: prudential_train=Prudential_train.drop(axis=1,labels=["Product_Info_2"])
```

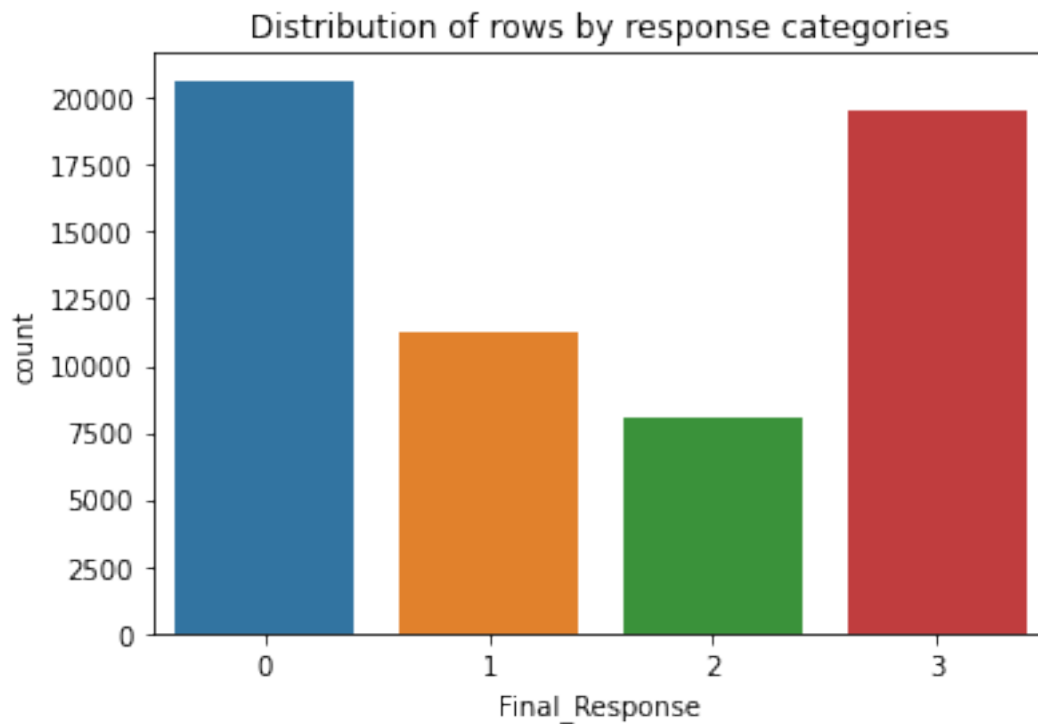
```
[21]: prudential_train.dropna(axis=1,inplace=True)
```

```
[22]: def new_target(row):
    if (row['Response']<=5):
        val=0
    elif (row['Response']==6):
        val=1
    elif (row['Response']==7):
        val=2
    elif (row['Response']==8):
        val=3

    else:
        val=-1
    return val
prudential_train['Final_Response']=prudential_train.apply(new_target,axis=1)
```

```
[23]: sns.countplot(x=prudential_train.Final_Response).set_title('Distribution of_
    ↳rows by response categories')
```

```
[23]: Text(0.5, 1.0, 'Distribution of rows by response categories')
```



10 Base Model

```
[24]: y = prudential_train.Final_Response
X = prudential_train.drop(labels=['Response'],axis=1)
X = X.drop(labels=['Final_Response'],axis=1)
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=.
↪20,random_state=1)
```

```
print("Shape of X_train dataset {}".format(X_train.shape))
print("Shape of X_test dataset {}".format(X_test.shape))

print("Shape of y_train dataset {}".format(y_train.shape))
print("Shape of y_valid dataset {}".format(y_test.shape))
```

```
Shape of X_train dataset (47504, 113)
Shape of X_test dataset (11877, 113)
Shape of y_train dataset (47504,)
Shape of y_valid dataset (11877,)
```

11 Decision Tree

```
[25]: model = DecisionTreeClassifier()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.514018691588785

```
=====
              precision    recall  f1-score   support

    0           0.60       0.58       0.59       4205
    1           0.32       0.34       0.33       2206
    2           0.30       0.31       0.30       1608
    3           0.64       0.62       0.63       3858

 accuracy                   0.51       11877
 macro avg           0.46       0.46       0.46       11877
weighted avg           0.52       0.51       0.52       11877
```

12 Naive Bayes

```
[26]: model = GaussianNB()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.428222615138503

```
=====
              precision    recall  f1-score   support

    0           0.71       0.21       0.33       4205
    1           0.26       0.07       0.11       2206
    2           0.29       0.22       0.25       1608
    3           0.42       0.95       0.58       3858

 accuracy                   0.43       11877
 macro avg           0.42       0.36       0.32       11877
weighted avg           0.47       0.43       0.36       11877
```


13 Treating Outliers

```
[29]: def remove_outlier(df_in, col_name):
      q1 = df_in[col_name].quantile(0.25)
      q3 = df_in[col_name].quantile(0.75)
      iqr = q3-q1 #Interquartile range
      fence_low = q1-1.5*iqr
      fence_high = q3+1.5*iqr
      df_out = df_in.loc[(df_in[col_name] > fence_low) & (df_in[col_name] <=
      ↪fence_high)]
      return df_out

dev=remove_outlier(Prudential_train,'BMI')
dev=remove_outlier(dev,'Wt')
dev=remove_outlier(dev,'Ht')
```

```
[30]: sns.boxplot(data=dev[misc_cols])
```

```
[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa03fe1da90>
```



```
[31]: prudential_X_train = dev
```

```
[32]: def new_target(row):
      if (row['Response']<=5):
```

```

        val=0
    elif (row['Response']==6):
        val=1
    elif (row['Response']==7):
        val=2
    elif (row['Response']==8):
        val=3

    else:
        val=-1
    return val
prudential_X_train['Final_Response']=prudential_X_train.apply(new_target,axis=1)

```

```
[33]: medical_keyword_cols=[col for col in prudential_X_train.columns if str(col).
    ↳startswith("Medical_Keyword")]
```

```
[34]: medical_cols=[col for col in prudential_X_train.columns if str(col).
    ↳startswith("Medical_History")]
```

```
[35]: prudential_X_train['Total_MedKwrds']=prudential_X_train[medical_keyword_cols].
    ↳sum(axis=1)
prudential_X_train['Total_MedHist']=prudential_X_train[medical_cols].sum(axis=1)
```

```
[36]: prudential_X_train['Total_MedKwrds']
```

```
[36]: 0      0
      1      0
      2      0
      3      1
      4      0
      ..
59376    0
59377    0
59378    1
59379    2
59380    0
Name: Total_MedKwrds, Length: 57348, dtype: int64
```

```
[37]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
prudential_X_train['Product_Info_2_en'] = le.
    ↳fit_transform(prudential_X_train['Product_Info_2'])
```

```
[38]: prudential_X_train['Product_Info_2_en']
```

```
[38]: 0      16
      1      0
```

```

2         18
3         17
4         15
..
59376     14
59377     16
59378     18
59379     15
59380      7
Name: Product_Info_2_en, Length: 57348, dtype: int64

```

```
[39]: prudential_X_train = prudential_X_train.drop(axis=1, labels=['Product_Info_2'])
```

```
[40]: prudential_X_train.Final_Response.unique()
```

```
[40]: array([3, 0, 1, 2])
```

14 Feature Selection

```
[41]: prudential_X_train = prudential_X_train.drop(labels = ['Response'], axis=1)
prudential_X_train = prudential_X_train.drop(labels = 
↳ ['Medical_History_10'], axis=1)
prudential_X_train = prudential_X_train.drop(labels = ['Id'], axis=1)
```

15 Fill Null Values

```
[42]: prudential_X_train = prudential_X_train.fillna(prudential_X_train.mean())
```

16 Build Model

```
[43]: y = prudential_X_train.Final_Response
X = prudential_X_train.drop(labels=['Final_Response'], axis=1)
X = X.fillna(X.mean())

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.
↳ 20, random_state=1)
print("Shape of X_train dataset {}".format(X_train.shape))
print("Shape of X_test dataset {}".format(X_test.shape))

print("Shape of y_train dataset {}".format(y_train.shape))
print("Shape of y_valid dataset {}".format(y_test.shape))
```

Shape of X_train dataset (45878, 127)
 Shape of X_test dataset (11470, 127)
 Shape of y_train dataset (45878,)
 Shape of y_valid dataset (11470,)

17 Decision Tree - with feature selection

```
[44]: model = DecisionTreeClassifier()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.5431560592850916

```
=====
```

	precision	recall	f1-score	support
0	0.60	0.59	0.59	3731
1	0.39	0.42	0.40	2202
2	0.35	0.35	0.35	1642
3	0.68	0.65	0.66	3895
accuracy			0.54	11470
macro avg	0.50	0.50	0.50	11470
weighted avg	0.55	0.54	0.55	11470

18 Naive Bayes - with feature selection

```
[46]: model = GaussianNB()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.44533565823888405

```
=====
```

	precision	recall	f1-score	support
0	0.58	0.32	0.41	3731
1	0.27	0.10	0.14	2202
2	0.24	0.36	0.29	1642
3	0.50	0.80	0.62	3895

accuracy			0.45	11470
macro avg	0.40	0.39	0.37	11470
weighted avg	0.45	0.45	0.41	11470

[]: