

Decision Tree, Naive Bayes and Neural Network

November 13, 2021

```
[5]: import pandas as pd
import matplotlib.pyplot as plt
from collections import Counter
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB

from sklearn.metrics import roc_curve, auc
from sklearn import metrics
from sklearn.neural_network import MLPClassifier
```

```
[8]: Prudential_train = pd.read_csv("https://raw.githubusercontent.com/crisajose/
↳CIND-820-Big-Data-Analytics-Project/main/train.csv")
```

```
[9]: Prudential_train.head()
```

```
[9]:
```

	Id	Product_Info_1	Product_Info_2	Product_Info_3	Product_Info_4	\
0	2	1	D3	10	0.076923	
1	5	1	A1	26	0.076923	
2	6	1	E1	26	0.076923	
3	7	1	D4	10	0.487179	
4	8	1	D2	26	0.230769	

	Product_Info_5	Product_Info_6	Product_Info_7	Ins_Age	Ht	...	\
0	2	1	1	0.641791	0.581818	...	
1	2	3	1	0.059701	0.600000	...	
2	2	3	1	0.029851	0.745455	...	
3	2	3	1	0.164179	0.672727	...	
4	2	3	1	0.417910	0.654545	...	

	Medical_Keyword_40	Medical_Keyword_41	Medical_Keyword_42	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	

4	0	0	0
---	---	---	---

	Medical_Keyword_43	Medical_Keyword_44	Medical_Keyword_45	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	Medical_Keyword_46	Medical_Keyword_47	Medical_Keyword_48	Response
0	0	0	0	8
1	0	0	0	4
2	0	0	0	8
3	0	0	0	8
4	0	0	0	8

[5 rows x 128 columns]

```
[10]: CATEGORICAL_COLUMNS = ["Product_Info_1", "Product_Info_2", "Product_Info_3",\
    ↪ "Product_Info_5", "Product_Info_6",\
    ↪ "Product_Info_7", "Employment_Info_2",\
    ↪ "Employment_Info_3", "Employment_Info_5", "InsuredInfo_1",\
    ↪ "InsuredInfo_2", "InsuredInfo_3", "InsuredInfo_4",\
    ↪ "InsuredInfo_5", "InsuredInfo_6", "InsuredInfo_7",\
    ↪ "Insurance_History_1", "Insurance_History_2",\
    ↪ "Insurance_History_3", "Insurance_History_4", "Insurance_History_7",\
    ↪ "Insurance_History_8", "Insurance_History_9",\
    ↪ "Family_Hist_1", "Medical_History_2", "Medical_History_3",\
    ↪ "Medical_History_4", "Medical_History_5",\
    ↪ "Medical_History_6", "Medical_History_7", "Medical_History_8",\
    ↪ "Medical_History_9", "Medical_History_11",\
    ↪ "Medical_History_12", "Medical_History_13", "Medical_History_14",\
    ↪ "Medical_History_16", "Medical_History_17",\
    ↪ "Medical_History_18", "Medical_History_19", "Medical_History_20",\
    ↪ "Medical_History_21", "Medical_History_22",\
    ↪ "Medical_History_23", "Medical_History_25", "Medical_History_26",\
    ↪ "Medical_History_27", "Medical_History_28",\
    ↪ "Medical_History_29", "Medical_History_30", "Medical_History_31",\
    ↪ "Medical_History_33", "Medical_History_34",\
    ↪ "Medical_History_35", "Medical_History_36", "Medical_History_37",\
    ↪ "Medical_History_38", "Medical_History_39",\
    ↪ "Medical_History_40", "Medical_History_41"]
CONTINUOUS_COLUMNS = ["Product_Info_4", "Ins_Age", "Ht", "Wt", "BMI",
    ↪ "Employment_Info_1", "Employment_Info_4",\
    ↪ "Employment_Info_6",
```

```

        "Insurance_History_5", "Family_Hist_2", "Family_Hist_3",
        ↪ "Family_Hist_4", "Family_Hist_5"]
DISCRETE_COLUMNS = ["Medical_History_1", "Medical_History_10",
        ↪ "Medical_History_15", "Medical_History_24", "Medical_History_32"]
DUMMY_COLUMNS = ["Medical_Keyword_{}".format(i) for i in range(1, 48)]

```

```
[11]: categorical_data = Prudential_train[CATEGORICAL_COLUMNS]
```

```
[12]: continuous_data = Prudential_train[CONTINUOUS_COLUMNS]
```

```
[13]: discrete_data = Prudential_train[DISCRETE_COLUMNS]
```

```
[14]: dummy_data = Prudential_train[DUMMY_COLUMNS]
```

1 Variable Types

```
[16]: Prudential_train.dtypes
```

```

[16]: Id                int64
      Product_Info_1    int64
      Product_Info_2    object
      Product_Info_3    int64
      Product_Info_4    float64
      ...
      Medical_Keyword_45 int64
      Medical_Keyword_46 int64
      Medical_Keyword_47 int64
      Medical_Keyword_48 int64
      Response          int64
      Length: 128, dtype: object

```

2 NULL values

```

[17]: Prudential_NULL = Prudential_train.isnull().any()
      Prudential_NULL = [col for col in Prudential_train.columns if
        ↪ Prudential_train[col].isnull().any()]
      Prudential_NULL

```

```

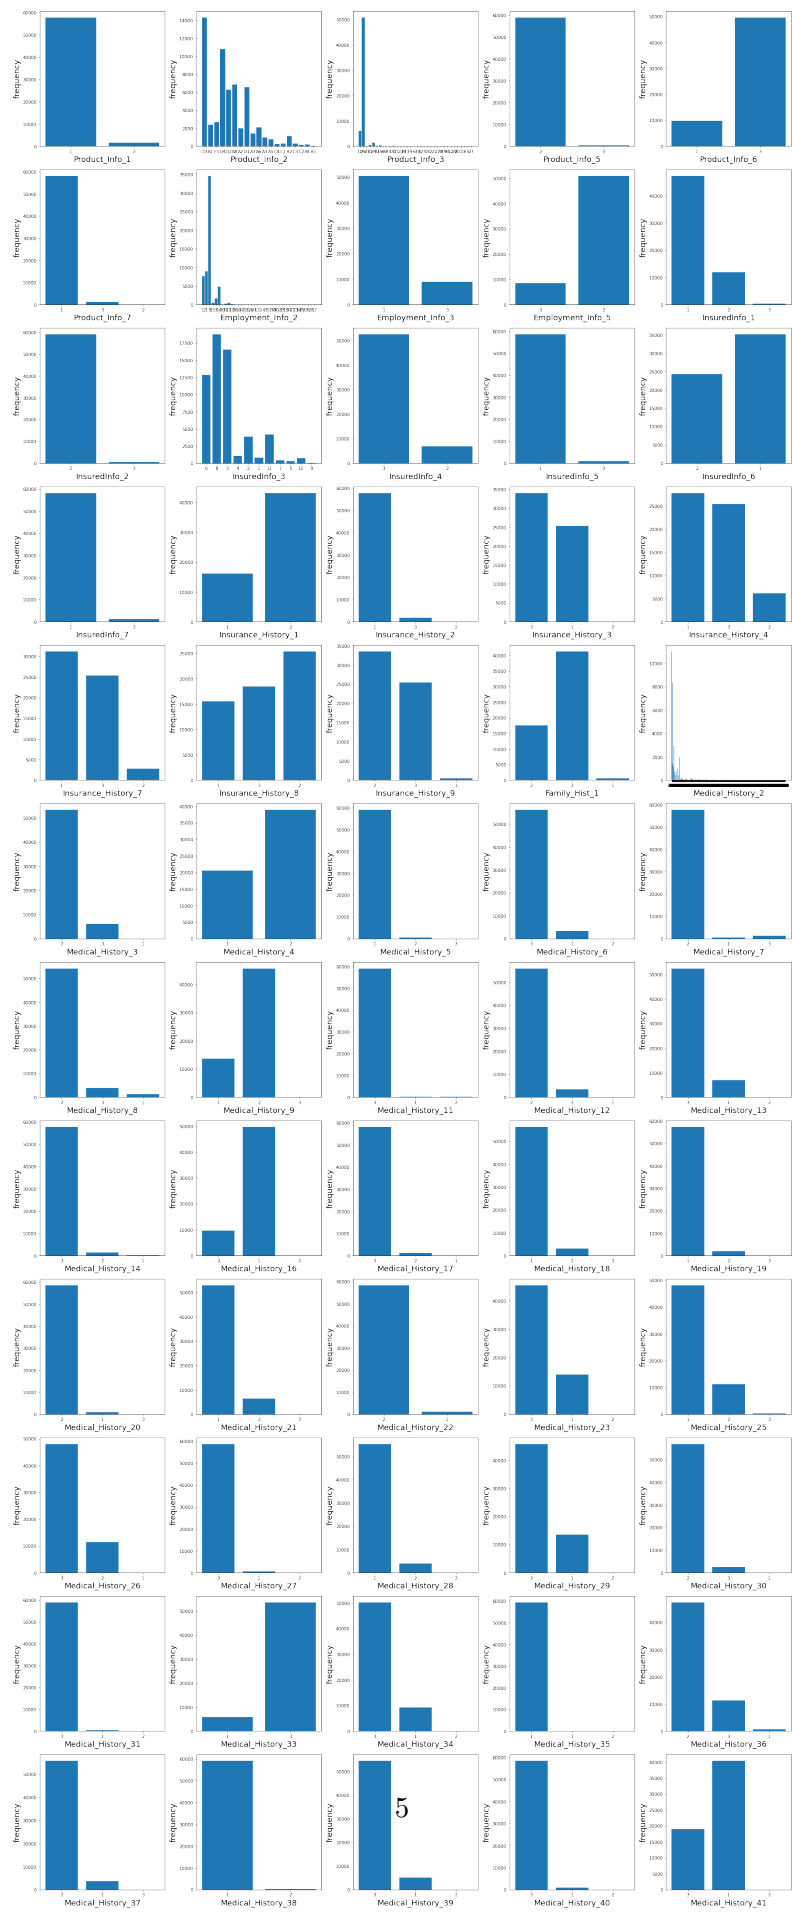
[17]: ['Employment_Info_1',
      'Employment_Info_4',
      'Employment_Info_6',
      'Insurance_History_5',
      'Family_Hist_2',

```

```
'Family_Hist_3',  
'Family_Hist_4',  
'Family_Hist_5',  
'Medical_History_1',  
'Medical_History_10',  
'Medical_History_15',  
'Medical_History_24',  
'Medical_History_32']
```

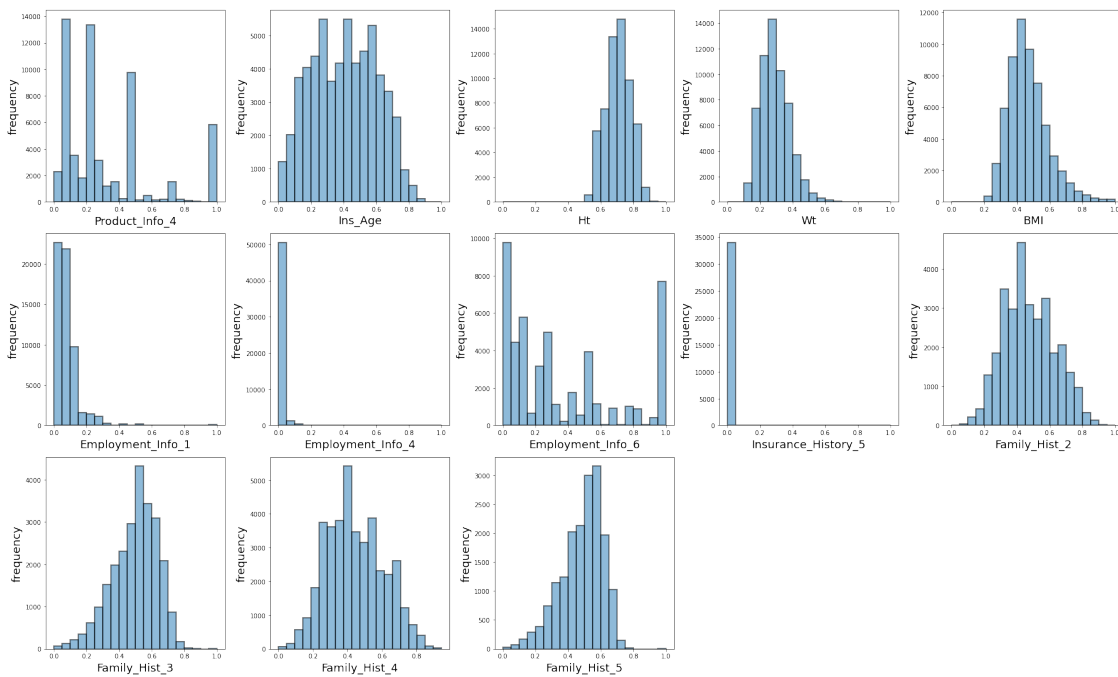
3 Categorical Variable - Plot

```
[18]: def plot_categoricals(data):  
    ncols = len(data.columns)  
    fig = plt.figure(figsize=(5 * 5, 5 * (ncols // 5 + 1)))  
    for i, col in enumerate(data.columns):  
        cnt = Counter(data[col])  
        keys = list(cnt.keys())  
        vals = list(cnt.values())  
        plt.subplot(ncols // 5 + 1, 5, i + 1)  
        plt.bar(range(len(keys)), vals, align="center")  
        plt.xticks(range(len(keys)), keys)  
        plt.xlabel(col, fontsize=18)  
        plt.ylabel("frequency", fontsize=18)  
    fig.tight_layout()  
    plt.show()  
  
plot_categoricals(categorical_data)
```




4 Continuous Variable - Plot

```
[19]: def plot_histgrams(data):
    ncols = len(data.columns)
    fig = plt.figure(figsize=(5 * 5, 5 * (ncols // 5 + 1)))
    for i, col in enumerate(data.columns):
        X = data[col].dropna()
        plt.subplot(ncols // 5 + 1, 5, i + 1)
        plt.hist(X, bins=20, alpha=0.5, \
                 edgecolor="black", linewidth=2.0)
        plt.xlabel(col, fontsize=18)
        plt.ylabel("frequency", fontsize=18)
    fig.tight_layout()
    plt.show()
plot_histgrams(continuous_data)
```



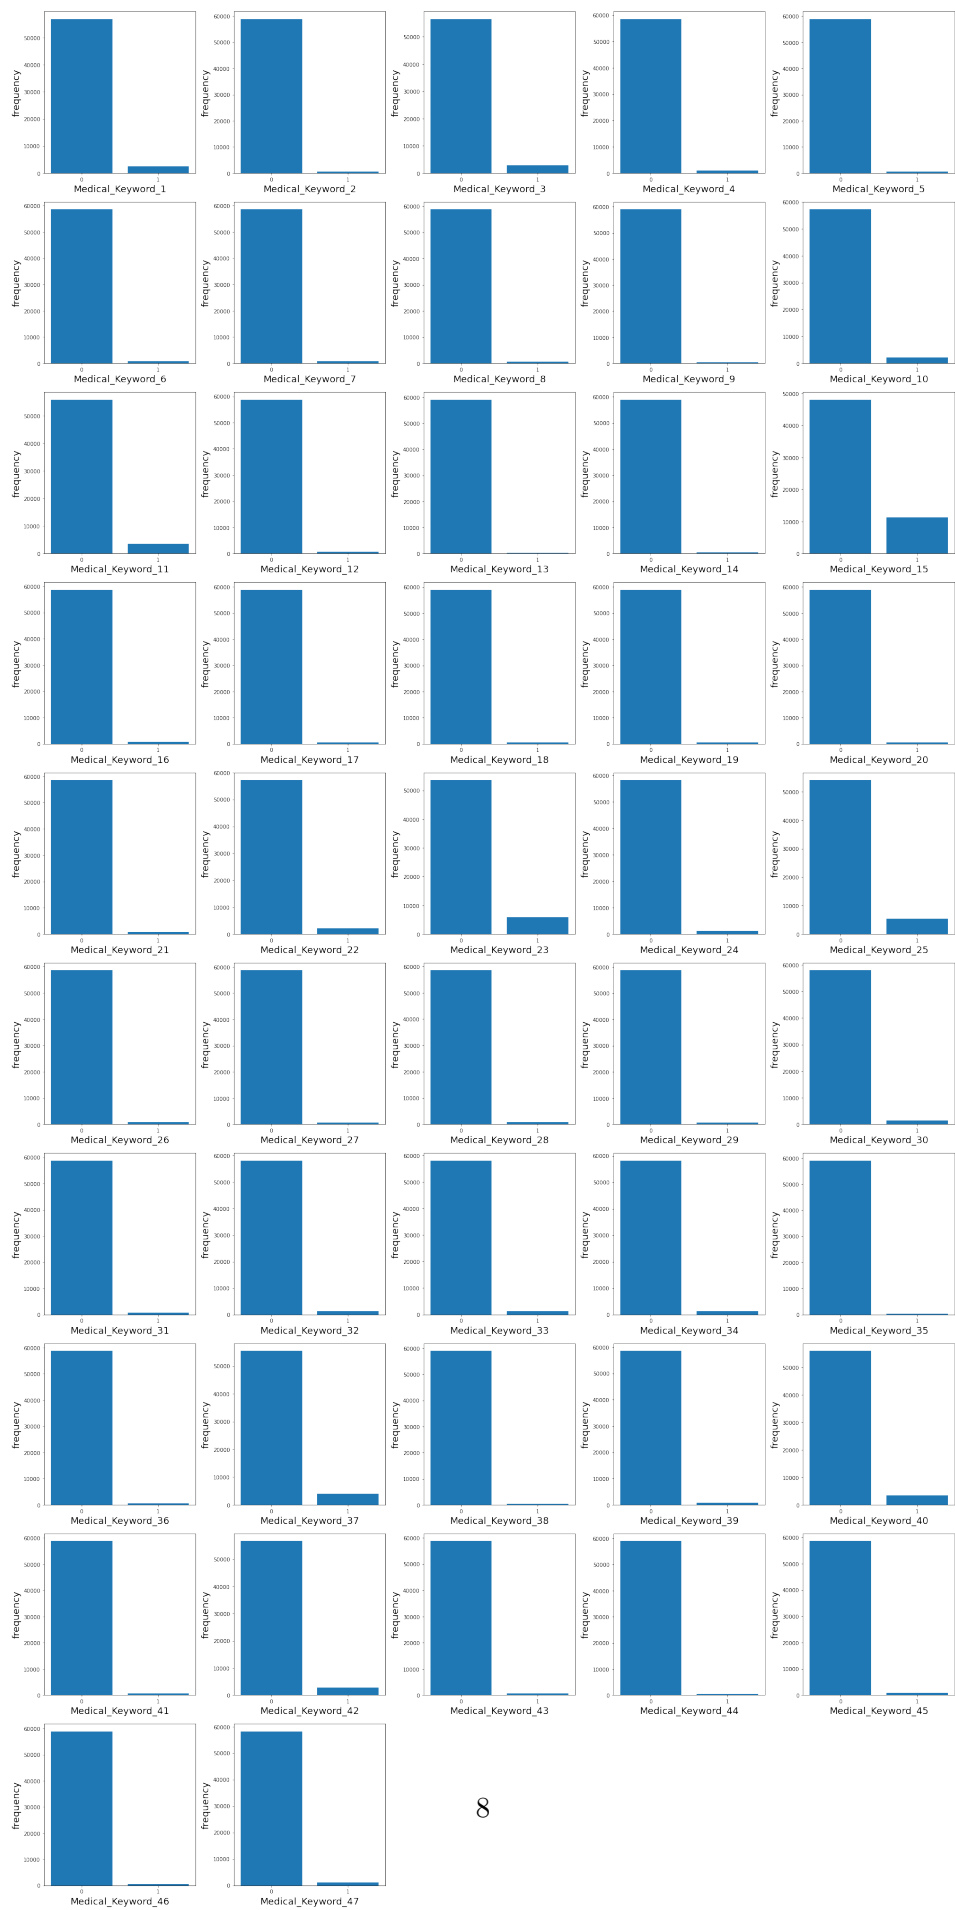
5 Discrete Variable - Plot

```
[20]: plot_histograms(discrete_data)
```



6 Dummy Variable - Plot


```
[21]: plot_categoricals(dummy_data)
```



7 Response Data Distribution

```
[22]: sns.countplot(x=Prudential_train.Response).set_title('Distribution of rows by_  
→response categories')
```


```
[22]: Text(0.5, 1.0, 'Distribution of rows by response categories')
```



8 Outliers Plot

```
[23]: misc_cols=["Ins_Age", "Ht", "Wt", "BMI"]  
sns.boxplot(data=Prudential_train[misc_cols])
```

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7f750b12db90>
```



9 Reassign Risk Class

```
[24]: prudential_train=Prudential_train.drop(axis=1,labels=["Product_Info_2"])
```


```
[25]: prudential_train.dropna(axis=1,inplace=True)
```

```
[26]: def new_target(row):
    if (row['Response']<=5):
        val=0
    elif (row['Response']==6):
        val=1
    elif (row['Response']==7):
        val=2
    elif (row['Response']==8):
        val=3

    else:
        val=-1
    return val
prudential_train['Final_Response']=prudential_train.apply(new_target,axis=1)
```

```
[27]: sns.countplot(x=prudential_train.Final_Response).set_title('Distribution of_
    ↳rows by response categories')
```

```
[27]: Text(0.5, 1.0, 'Distribution of rows by response categories')
```



10 Base Model

```
[28]: y = prudential_train.Final_Response
X = prudential_train.drop(labels=['Response'],axis=1)
X = X.drop(labels=['Final_Response'],axis=1)
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=.
↪20,random_state=1)
```

```
print("Shape of X_train dataset {}".format(X_train.shape))
print("Shape of X_test dataset {}".format(X_test.shape))

print("Shape of y_train dataset {}".format(y_train.shape))
print("Shape of y_valid dataset {}".format(y_test.shape))
```

```
Shape of X_train dataset (47504, 113)
Shape of X_test dataset (11877, 113)
Shape of y_train dataset (47504,)
Shape of y_valid dataset (11877,)
```

11 Decision Tree

```
[29]: model = DecisionTreeClassifier()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.5117453902500632

```
=====
              precision    recall  f1-score   support

    0           0.60       0.58       0.59       4205
    1           0.32       0.33       0.32       2206
    2           0.28       0.31       0.30       1608
    3           0.64       0.62       0.63       3858

 accuracy                   0.51       11877
 macro avg           0.46       0.46       0.46       11877
weighted avg           0.52       0.51       0.51       11877
```

12 Naive Bayes

```
[30]: model = GaussianNB()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.428222615138503

```
=====
              precision    recall  f1-score   support

    0           0.71       0.21       0.33       4205
    1           0.26       0.07       0.11       2206
    2           0.29       0.22       0.25       1608
    3           0.42       0.95       0.58       3858

 accuracy                   0.43       11877
 macro avg           0.42       0.36       0.32       11877
weighted avg           0.47       0.43       0.36       11877
```


13 Treating Outliers

```
[31]: def remove_outlier(df_in, col_name):
      q1 = df_in[col_name].quantile(0.25)
      q3 = df_in[col_name].quantile(0.75)
      iqr = q3-q1 #Interquartile range
      fence_low = q1-1.5*iqr
      fence_high = q3+1.5*iqr
      df_out = df_in.loc[(df_in[col_name] > fence_low) & (df_in[col_name] <=
      ↪fence_high)]
      return df_out

dev=remove_outlier(Prudential_train,'BMI')
dev=remove_outlier(dev,'Wt')
dev=remove_outlier(dev,'Ht')
```

```
[32]: sns.boxplot(data=dev[misc_cols])
```

```
[32]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7503022950>
```



```
[33]: prudential_X_train = dev
```

```
[34]: def new_target(row):
      if (row['Response']<=5):
```

```

        val=0
    elif (row['Response']==6):
        val=1
    elif (row['Response']==7):
        val=2
    elif (row['Response']==8):
        val=3

    else:
        val=-1
    return val
prudential_X_train['Final_Response']=prudential_X_train.apply(new_target,axis=1)

```

```
[35]: medical_keyword_cols=[col for col in prudential_X_train.columns if str(col).
    ↳startswith("Medical_Keyword")]
```

```
[36]: medical_cols=[col for col in prudential_X_train.columns if str(col).
    ↳startswith("Medical_History")]
```

```
[37]: prudential_X_train['Total_MedKwrds']=prudential_X_train[medical_keyword_cols].
    ↳sum(axis=1)
prudential_X_train['Total_MedHist']=prudential_X_train[medical_cols].sum(axis=1)
```

```
[39]: prudential_X_train['Total_MedKwrds']
```

```
[39]: 0      0
      1      0
      2      0
      3      1
      4      0
      ..
59376    0
59377    0
59378    1
59379    2
59380    0
Name: Total_MedKwrds, Length: 57348, dtype: int64
```

```
[40]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
prudential_X_train['Product_Info_2_en'] = le.
    ↳fit_transform(prudential_X_train['Product_Info_2'])
```

```
[41]: prudential_X_train['Product_Info_2_en']
```

```
[41]: 0      16
      1      0
```

```

2      18
3      17
4      15
..
59376   14
59377   16
59378   18
59379   15
59380    7
Name: Product_Info_2_en, Length: 57348, dtype: int64

```

```
[42]: prudential_X_train = prudential_X_train.drop(axis=1, labels=['Product_Info_2'])
```

```
[43]: prudential_X_train.Final_Response.unique()
```

```
[43]: array([3, 0, 1, 2])
```

14 Feature Selection

```
[44]: prudential_X_train = prudential_X_train.drop(labels = ['Response'], axis=1)
prudential_X_train = prudential_X_train.drop(labels = ['Medical_History_10'], axis=1)
prudential_X_train = prudential_X_train.drop(labels = ['Id'], axis=1)
```

15 Fill Null Values

```
[45]: prudential_X_train = prudential_X_train.fillna(prudential_X_train.mean())
```

16 Build Model

```
[46]: y = prudential_X_train.Final_Response
X = prudential_X_train.drop(labels=['Final_Response'], axis=1)
X = X.fillna(X.mean())

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.
    ↳ 20, random_state=1)
print("Shape of X_train dataset {}".format(X_train.shape))
print("Shape of X_test dataset {}".format(X_test.shape))

print("Shape of y_train dataset {}".format(y_train.shape))
print("Shape of y_valid dataset {}".format(y_test.shape))
```

Shape of X_train dataset (45878, 127)
 Shape of X_test dataset (11470, 127)
 Shape of y_train dataset (45878,)
 Shape of y_valid dataset (11470,)

17 Decision Tree - with feature selection

```
[47]: model = DecisionTreeClassifier()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.542458587619878

```
=====
```

	precision	recall	f1-score	support
0	0.60	0.59	0.59	3731
1	0.39	0.41	0.40	2202
2	0.33	0.34	0.33	1642
3	0.68	0.66	0.67	3895
accuracy			0.54	11470
macro avg	0.50	0.50	0.50	11470
weighted avg	0.55	0.54	0.54	11470

18 Naive Bayes - with feature selection

```
[48]: model = GaussianNB()
model.fit(X_train, y_train)
model_predictions = model.predict(X_test)
print("Accuracy score: {}".format(accuracy_score(y_test, model_predictions)))
print("="*80)
print(classification_report(y_test, model_predictions))
```

Accuracy score: 0.44533565823888405

```
=====
```

	precision	recall	f1-score	support
0	0.58	0.32	0.41	3731
1	0.27	0.10	0.14	2202
2	0.24	0.36	0.29	1642
3	0.50	0.80	0.62	3895

accuracy			0.45	11470
macro avg	0.40	0.39	0.37	11470
weighted avg	0.45	0.45	0.41	11470

19 Deep Neural Network

```
[2]: !pip install tensorflow
```

```
Collecting tensorflow
  Using cached tensorflow-2.7.0-cp37-cp37m-manylinux2010_x86_64.whl (489.6 MB)
Requirement already satisfied: six>=1.12.0 in /opt/conda/lib/python3.7/site-
packages (from tensorflow) (1.14.0)
Collecting tensorboard~=2.6
  Using cached tensorboard-2.7.0-py3-none-any.whl (5.8 MB)
Requirement already satisfied: wheel<1.0,>=0.32.0 in
/opt/conda/lib/python3.7/site-packages (from tensorflow) (0.34.2)
Collecting grpcio<2.0,>=1.24.3
  Using cached
grpcio-1.41.1-cp37-cp37m-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.9 MB)
Collecting tensorflow-io-gcs-filesystem>=0.21.0
  Using cached tensorflow_io_gcs_filesystem-0.22.0-cp37-cp37m-manylinux_2_12_x86
_64.manylinux2010_x86_64.whl (2.1 MB)
Collecting gast<0.5.0,>=0.2.1
  Using cached gast-0.4.0-py3-none-any.whl (9.8 kB)
Collecting absl-py>=0.4.0
  Using cached absl_py-1.0.0-py3-none-any.whl (126 kB)
Processing ./cache/pip/wheels/3f/e3/ec/8a8336ff196023622fbc36de0c5a5c218cbb241
11d1d4c7f2/termcolor-1.1.0-py3-none-any.whl
Collecting opt-einsum>=2.3.2
  Using cached opt_einsum-3.3.0-py3-none-any.whl (65 kB)
Requirement already satisfied: typing-extensions>=3.6.6 in
/opt/conda/lib/python3.7/site-packages (from tensorflow) (3.7.4.2)
Requirement already satisfied: protobuf>=3.9.2 in /opt/conda/lib/python3.7/site-
packages (from tensorflow) (3.11.4)
Collecting astunparse>=1.6.0
  Using cached astunparse-1.6.3-py2.py3-none-any.whl (12 kB)
Collecting flatbuffers<3.0,>=1.12
  Using cached flatbuffers-2.0-py2.py3-none-any.whl (26 kB)
Collecting wrapt>=1.11.0
  Using cached wrapt-1.13.3-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.ma
nylinux_2_12_x86_64.manylinux2010_x86_64.whl (79 kB)
Collecting tensorflow-estimator<2.8,~=2.7.0rc0
  Using cached tensorflow_estimator-2.7.0-py2.py3-none-any.whl (463 kB)
Requirement already satisfied: h5py>=2.9.0 in /opt/conda/lib/python3.7/site-
```

```

packages (from tensorflow) (2.10.0)
Collecting libclang>=9.0.1
  Using cached libclang-12.0.0-py2.py3-none-manylinux1_x86_64.whl (13.4 MB)
Collecting keras-preprocessing>=1.1.1
  Using cached Keras_Preprocessing-1.1.2-py2.py3-none-any.whl (42 kB)
Collecting keras<2.8,>=2.7.0rc0
  Using cached keras-2.7.0-py2.py3-none-any.whl (1.3 MB)
Collecting google-pasta>=0.1.1
  Using cached google_pasta-0.2.0-py3-none-any.whl (57 kB)
Requirement already satisfied: numpy>=1.14.5 in /opt/conda/lib/python3.7/site-
packages (from tensorflow) (1.18.4)
Collecting werkzeug>=0.11.15
  Using cached Werkzeug-2.0.2-py3-none-any.whl (288 kB)
Collecting markdown>=2.6.8
  Using cached Markdown-3.3.4-py3-none-any.whl (97 kB)
Requirement already satisfied: google-auth<3,>=1.6.3 in
/opt/conda/lib/python3.7/site-packages (from tensorboard~=2.6->tensorflow)
(1.16.1)
Requirement already satisfied: requests<3,>=2.21.0 in
/opt/conda/lib/python3.7/site-packages (from tensorboard~=2.6->tensorflow)
(2.23.0)
Collecting google-auth-oauthlib<0.5,>=0.4.1
  Using cached google_auth_oauthlib-0.4.6-py2.py3-none-any.whl (18 kB)
Collecting tensorboard-plugin-wit>=1.6.0
  Using cached tensorboard_plugin_wit-1.8.0-py3-none-any.whl (781 kB)
Requirement already satisfied: setuptools>=41.0.0 in
/opt/conda/lib/python3.7/site-packages (from tensorboard~=2.6->tensorflow)
(46.1.3.post20200325)
Collecting tensorboard-data-server<0.7.0,>=0.6.0
  Using cached tensorboard_data_server-0.6.1-py3-none-manylinux2010_x86_64.whl
(4.9 MB)
Requirement already satisfied: importlib-metadata; python_version < "3.8" in
/opt/conda/lib/python3.7/site-packages (from
markdown>=2.6.8->tensorboard~=2.6->tensorflow) (1.6.0)
Requirement already satisfied: pyasn1-modules>=0.2.1 in
/opt/conda/lib/python3.7/site-packages (from google-
auth<3,>=1.6.3->tensorboard~=2.6->tensorflow) (0.2.8)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in
/opt/conda/lib/python3.7/site-packages (from google-
auth<3,>=1.6.3->tensorboard~=2.6->tensorflow) (4.1.0)
Requirement already satisfied: rsa<4.1,>=3.1.4 in /opt/conda/lib/python3.7/site-
packages (from google-auth<3,>=1.6.3->tensorboard~=2.6->tensorflow) (4.0)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in
/opt/conda/lib/python3.7/site-packages (from
requests<3,>=2.21.0->tensorboard~=2.6->tensorflow) (1.25.9)
Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-
packages (from requests<3,>=2.21.0->tensorboard~=2.6->tensorflow) (2.9)
Requirement already satisfied: certifi>=2017.4.17 in

```

```

/opt/conda/lib/python3.7/site-packages (from
requests<3,>=2.21.0->tensorboard~=2.6->tensorflow) (2020.4.5.2)
Requirement already satisfied: chardet<4,>=3.0.2 in
/opt/conda/lib/python3.7/site-packages (from
requests<3,>=2.21.0->tensorboard~=2.6->tensorflow) (3.0.4)
Requirement already satisfied: requests-oauthlib>=0.7.0 in
/opt/conda/lib/python3.7/site-packages (from google-auth-
oauthlib<0.5,>=0.4.1->tensorboard~=2.6->tensorflow) (1.3.0)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-
packages (from importlib-metadata; python_version <
"3.8"->markdown>=2.6.8->tensorboard~=2.6->tensorflow) (3.1.0)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in
/opt/conda/lib/python3.7/site-packages (from pyasn1-modules>=0.2.1->google-
auth<3,>=1.6.3->tensorboard~=2.6->tensorflow) (0.4.8)
Requirement already satisfied: oauthlib>=3.0.0 in /opt/conda/lib/python3.7/site-
packages (from requests-oauthlib>=0.7.0->google-auth-
oauthlib<0.5,>=0.4.1->tensorboard~=2.6->tensorflow) (3.0.1)
Installing collected packages: werkzeug, absl-py, markdown, grpcio, google-auth-
oauthlib, tensorboard-plugin-wit, tensorboard-data-server, tensorboard,
tensorflow-io-gcs-filesystem, gast, termcolor, opt-einsum, astunparse,
flatbuffers, wrapt, tensorflow-estimator, libclang, keras-preprocessing, keras,
google-pasta, tensorflow
Successfully installed absl-py-1.0.0 astunparse-1.6.3 flatbuffers-2.0 gast-0.4.0
google-auth-oauthlib-0.4.6 google-pasta-0.2.0 grpcio-1.41.1 keras-2.7.0 keras-
preprocessing-1.1.2 libclang-12.0.0 markdown-3.3.4 opt-einsum-3.3.0
tensorboard-2.7.0 tensorboard-data-server-0.6.1 tensorboard-plugin-wit-1.8.0
tensorflow-2.7.0 tensorflow-estimator-2.7.0 tensorflow-io-gcs-filesystem-0.22.0
termcolor-1.1.0 werkzeug-2.0.2 wrapt-1.13.3

```

```

[51]: from tensorflow import keras
def get_model():
    model = keras.Sequential([
        keras.layers.Flatten(input_shape=[X.shape[-1]]),
        keras.layers.Dense(512, activation='relu'),
        keras.layers.Dense(256, activation='relu'),
        keras.layers.Dense(128, activation='relu'),
        keras.layers.Dense(64, activation='relu'),
        keras.layers.Dense(32, activation='relu'),
        keras.layers.Dropout(0.5),
        keras.layers.Dense(9, activation='softmax')
    ])
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    return model

```

```
[52]: # DNN (fit and validation)
import tensorflow as tf

batch_size = 512
train_ds = tf.data.Dataset.from_tensor_slices((X_train.values, y_train.values)).
    ↪ shuffle(len(X_train)).batch(batch_size)
val_ds = tf.data.Dataset.from_tensor_slices((X_test.values, y_test.values)).
    ↪ batch(batch_size)

model = get_model()
fit = model.fit(train_ds, validation_data=val_ds, epochs=20)
```

Epoch 1/20

90/90 [=====] - 2s 13ms/step - loss: 2.1382 - accuracy: 0.3099 - val_loss: 1.4310 - val_accuracy: 0.4064

Epoch 2/20

90/90 [=====] - 1s 11ms/step - loss: 1.5067 - accuracy: 0.3535 - val_loss: 1.2871 - val_accuracy: 0.4604

Epoch 3/20

90/90 [=====] - 1s 10ms/step - loss: 1.4167 - accuracy: 0.3716 - val_loss: 1.2783 - val_accuracy: 0.4879

Epoch 4/20

90/90 [=====] - 1s 10ms/step - loss: 1.3693 - accuracy: 0.3770 - val_loss: 1.2494 - val_accuracy: 0.4350

Epoch 5/20

90/90 [=====] - 1s 11ms/step - loss: 1.3209 - accuracy: 0.4000 - val_loss: 1.2194 - val_accuracy: 0.4800

Epoch 6/20

90/90 [=====] - 1s 10ms/step - loss: 1.2899 - accuracy: 0.4305 - val_loss: 1.2015 - val_accuracy: 0.4934

Epoch 7/20

90/90 [=====] - 1s 10ms/step - loss: 1.2658 - accuracy: 0.4559 - val_loss: 1.1972 - val_accuracy: 0.4903

Epoch 8/20

90/90 [=====] - 1s 10ms/step - loss: 1.2517 - accuracy: 0.4644 - val_loss: 1.1863 - val_accuracy: 0.5114

Epoch 9/20

90/90 [=====] - 1s 10ms/step - loss: 1.2370 - accuracy: 0.4733 - val_loss: 1.2382 - val_accuracy: 0.4317

Epoch 10/20

90/90 [=====] - 1s 10ms/step - loss: 1.2350 - accuracy: 0.4749 - val_loss: 1.1683 - val_accuracy: 0.4933

Epoch 11/20

90/90 [=====] - 1s 10ms/step - loss: 1.2046 - accuracy: 0.4920 - val_loss: 1.1334 - val_accuracy: 0.5140

Epoch 12/20

90/90 [=====] - 1s 10ms/step - loss: 1.2119 - accuracy:

```

0.4896 - val_loss: 1.1680 - val_accuracy: 0.4879
Epoch 13/20
90/90 [=====] - 1s 10ms/step - loss: 1.1981 - accuracy:
0.4973 - val_loss: 1.1245 - val_accuracy: 0.5176
Epoch 14/20
90/90 [=====] - 1s 10ms/step - loss: 1.1842 - accuracy:
0.5034 - val_loss: 1.1073 - val_accuracy: 0.5227
Epoch 15/20
90/90 [=====] - 1s 10ms/step - loss: 1.1688 - accuracy:
0.5090 - val_loss: 1.1394 - val_accuracy: 0.5159
Epoch 16/20
90/90 [=====] - 1s 10ms/step - loss: 1.1695 - accuracy:
0.5080 - val_loss: 1.0976 - val_accuracy: 0.5320
Epoch 17/20
90/90 [=====] - 1s 10ms/step - loss: 1.1567 - accuracy:
0.5117 - val_loss: 1.0920 - val_accuracy: 0.5303
Epoch 18/20
90/90 [=====] - 1s 11ms/step - loss: 1.1517 - accuracy:
0.5165 - val_loss: 1.1308 - val_accuracy: 0.5336
Epoch 19/20
90/90 [=====] - 1s 11ms/step - loss: 1.1430 - accuracy:
0.5204 - val_loss: 1.1563 - val_accuracy: 0.4853
Epoch 20/20
90/90 [=====] - 1s 11ms/step - loss: 1.1405 - accuracy:
0.5213 - val_loss: 1.0746 - val_accuracy: 0.5318

```

```

[53]: # DNN (accuracy by epoch)
import matplotlib.pyplot as plt

plt.plot(fit.history['accuracy'])
plt.plot(fit.history['val_accuracy'])
plt.show()

```

